

Bit Manipulation Templates

Quick Reference Table

#	Name	Description	Intuition	Time	Space	Key Implementation
136	Single Number	Array where every element appears twice except one. Find it.	XOR is its own inverse, so every duplicated pair cancels to 0 and only the lone element is left standing.	$O(n)$	$O(1)$	Standard
268	Missing Number	Array of n distinct numbers from $[0, n]$ with one missing. Find it.	XOR-ing every index against every value pairs each present number with its index — the missing number's index has no partner to cancel.	$O(n)$	$O(1)$	XOR each index i with <code>nums[i]</code> . Indices $0..n$ XOR'd with the n elements — the missing index has no pair.
137	Single Number II	Every element appears three times except one. Find it.	XOR cancels pairs (mod 2); here we need a counter mod 3, so two state bits (<code>ones</code> , <code>twos</code>) cycle each bit through $1 \rightarrow 2 \rightarrow 0$ and the survivor sits in <code>ones</code> .	$O(n)$	$O(1)$	Two-state XOR machine. <code>ones</code> tracks bits seen 1 mod 3 times, <code>twos</code> tracks bits seen 2 mod 3 times. When a bit reaches 3, it clears from both.
260	Single Number III	Two elements appear exactly once; all others appear twice. Find both.	XOR of all leaves $a \oplus b$; any set bit in it is a bit where a and b differ, so it splits the array into two groups each holding one unique.	$O(n)$	$O(1)$	XOR all $\rightarrow$ get $a \oplus b$. Use lowest set bit of $a \oplus b$ to partition nums into two groups (one containing a , one containing b). XOR each group independently.
191	Number of 1 Bits	Return the number of set bits in an unsigned integer.	Each $n \& (n-1)$ erases exactly one set bit, so the loop runs once per set bit — no need to scan all 32 positions.	$O(k)$ k = number of set bits	$O(1)$	Standard
338	Counting Bits	Return an array <code>result</code> where <code>result[i]</code> = number of 1 bits in i , for i in $[0, n]$.	i has the same set bits as $i \gg 1$ plus possibly its own last bit, so reuse the already-computed smaller answer instead of recounting.	$O(n)$	$O(n)$	DP relation — $i \gg 1$ is i with last bit removed (already computed), plus the last bit $i \& 1$.
190	Reverse Bits	Reverse the 32 bits of an unsigned integer.	Peel the lowest bit off n and stack it onto <code>result</code> from the bottom — after 32 shifts the bit order is fully mirrored.	$O(32) = O(1)$	$O(1)$	Fixed 32 iterations — each time take the LSB of n , append it to <code>result</code> , then shift both.
201	Bitwise AND of Numbers Range	Return the AND of all numbers in <code>[left, right]</code> .	AND only keeps bits that are 1 in every number; the low bits flip somewhere in the range, so only the common high-bit prefix of <code>left</code> and <code>right</code> survives.	$O(\log n)$	$O(1)$	right-shift both until equal to find shared prefix; shift back.
371	Sum of Two Integers	Return $a + b$ without using <code>+</code> or <code>-</code> .	Addition splits into a carry-free sum (XOR) and the carries (AND shifted left); feed the carry back until nothing carries over.	$O(32) = O(1)$	$O(1)$	XOR computes bit sum without carry; AND shifted left computes carry. Repeat until no carry.

#	Name	Description	Intuition	Time	Space	Key Implementation
318	Maximum Product of Word Lengths	Given a list of words, find the maximum product $\text{len}(a) * \text{len}(b)$ where a and b share no common letters.	Compress each word's letter set into a 26-bit integer so "share a letter?" becomes a single AND, replacing slow character-by-character comparison.	$O(n^2 + n \cdot k)$ $k = \text{avg word length}$	$O(n)$	encode each word as a 26-bit integer where bit $i = 1$ if letter 'a' + i appears. Two words share a letter iff their bitmasks have any common bit ($\text{mask}[i] \& \text{mask}[j] \neq 0$).
231	Power of Two	Given an integer n , return true if it is a power of two.	A power of two is a single 1 bit, and $n \& (n-1)$ clears that one bit to 0 — anything else leaves bits behind.	$O(1)$	$O(1)$	single-bit property of powers of 2
342	Power of Four	Given an integer n , return true if it is a power of four.	Powers of four are powers of two (one set bit) whose bit sits at an even position; masking against the odd-position mask $0xAAAAAAAA$ must give 0.	$O(1)$	$O(1)$	Power of four must be power of two (one set bit) AND that bit must be at an even bit position (0, 2, 4, 6, ...). Mask $0xAAAAAAAA$ has bits set at ALL odd positions; if $n \& 0xAAAAAAAA == 0$, the set bit is at an even position.
287	Find the Duplicate Number	Given an array of $n+1$ integers where each is in $[1, n]$, find the duplicate. No extra space, no modifying the array.	Following $i \rightarrow \text{nums}[i]$ turns the array into a linked list; a duplicate value means two indices point to the same node, forming a cycle whose entry is the duplicate.	$O(n)$	$O(1)$	Floyd's Cycle Detection. Treat the array as a linked list where $\text{nums}[i]$ is the next node. Since there's a duplicate, two indices point to the same value → creates a cycle. Find cycle entry = duplicate.
405	Convert a Number to Hexadecimal	Given an integer num , return its hexadecimal representation as a lowercase string. For negative numbers use two's complement.	Hex is base 16, so each group of 4 bits maps to one hex digit; mask the lowest 4 bits with $\text{num} \& 15$ and shift right 4 at a time, using an unsigned shift so the two's-complement sign bit fills with zeros.	$O(8) = O(1)$	$O(1)$	Standard
461	Hamming Distance	Given two integers x and y , return the Hamming distance — the number of bit positions where they differ.	XOR sets exactly the bits where x and y differ, so the answer is just the number of set bits in $x \wedge y$.	$O(k)$ $k = \text{differing bits}$	$O(1)$	Standard
476	Number Complement	Given a positive integer num , return its complement: flip every bit up to and including the most significant set bit.	Build a mask of all 1s spanning the bit width of num (from the MSB down to bit 0), then XOR — flipping exactly the meaningful bits while leaving the leading zeros untouched.	$O(\log \text{num})$	$O(1)$	Standard
957	Prison Cells After N Days	8 prison cells ($\text{cells}[i]$ is 0 or 1) update each day: a cell becomes 1 if both neighbors are equal, else 0. The two end cells always become 0 (they lack two neighbors). Return the state after n days.	With only 8 cells the state space is finite, so the configuration must cycle; detect the cycle length and reduce n modulo it to avoid simulating huge day counts.	$O(\min(n, \text{cycle}) \cdot 8)$	$O(\text{states})$	Standard

Canonical Template

Variables: `result` = XOR accumulator (lone survivor) · `n` = number being inspected/edited · `i` = bit position · `count` = set-bit counter · `mask` = 26-bit letter-presence set · masks: `1 << i` selects bit `i`, `n & (n-1)` clears lowest set bit, `n & (-n)` isolates lowest set bit **Pseudocode:**

XOR CANCEL:

```
result = 0
for each num in nums: result = result XOR num    (pairs cancel, lone survives)
```

BIT OPERATIONS on `n` at position `i`:

```
isSet = bit i of n is 1          -> (n >> i) & 1 == 1
set    = turn bit i on           -> n | (1 << i)
clear  = turn bit i off          -> n & ~(1 << i)
toggle = flip bit i              -> n ^ (1 << i)
lowest = isolate lowest set bit  -> n & (-n)
isPow2 = n > 0 and clearing lowest set bit gives 0
```

COUNT SET BITS (Brian Kernighan):

```
count = 0
while n != 0: count = count + 1; n = n & (n-1)    (each step removes one set bit)
```

BITMASK AS SET:

```
mask = 0
for each char c in word: mask = mask | (1 << (c - 'a'))
two words share no letter when (mask[i] & mask[j]) == 0
```

```
// bits are a set/counter – XOR cancels duplicates, AND/OR/shift edit individual bits. – WHEN: "appear"
// — XOR CANCEL — pairs cancel; lone element survives
```

```
int result = 0;
for (int num : nums) result ^= num;
```

```
// — BIT OPERATIONS —
```

```
boolean isSet = ((n >> i) & 1) == 1;    // check if bit i is set
int set       = n | (1 << i);           // set bit i
int clear     = n & ~(1 << i);           // clear bit i
int toggle    = n ^ (1 << i);           // toggle bit i
int lowest    = n & (-n);                // isolate lowest set bit
boolean isPow2 = n > 0 && (n & (n-1)) == 0;
```

```
// — COUNT SET BITS — Brian Kernighan: each iteration removes one set bit
```

```
// WHY n & (n-1) clears the lowest set bit:
```

```
//  n    = ...1000  (lowest set bit at this position)
//  n-1  = ...0111  (borrow flips that bit to 0 and all lower bits to 1)
//  n&(n-1)=...0000  (AND keeps higher bits, wipes the lowest set bit + below)
```

```
int count = 0;
while (n != 0) { count++; n &= (n - 1); } // n & (n-1) clears lowest set bit
```

```
// — BITMASK AS SET — 26-bit integer represents presence of 26 letters
```

```
int mask = 0;
for (char c : word.toCharArray()) mask |= (1 << (c - 'a'));
// check no overlap between two words:
if ((mask[i] & mask[j]) == 0) { /* no shared letter */ }
```

Variations

Variables: `ones / twos` = bits seen 1 / 2 times mod 3 · `xor` = `a ^ b` of the two uniques · `diff` = lowest differing bit (`xor & -xor`) · `a / b` = the two unique values · `dp[]` = set-bit counts by number · `shift` = bits dropped to reach common prefix · `carry` = AND-derived carry bits **Pseudocode:**

VARIATION 1 (mod-3 state machine, Single Number II):

```
ones = 0; twos = 0
for each num: ones = (ones XOR num) AND NOT twos; twos = (twos XOR num) AND NOT ones
after loop, ones holds the element appearing once
```

VARIATION 2 (split XOR into two groups, Single Number III):

```
xor = XOR of all nums           (= a XOR b)
diff = xor AND (-xor)           (lowest bit where a and b differ)
a = 0; for each num: if (num AND diff) != 0: a = a XOR num   (XOR one group)
b = xor XOR a                   (the other unique)
```

VARIATION 3 (DP relation, Counting Bits):

```
dp = array of size n+1
for i = 1..n: dp[i] = dp[i >> 1] + (i AND 1)   (drop last bit, add it back)
```

VARIATION 4 (common prefix, Bitwise AND of Range):

```
shift = 0
while left != right: left >>= 1; right >>= 1; shift = shift + 1
return left << shift           (shared high-bit prefix restored)
```

VARIATION 5 (carry loop, Sum of Two Integers):

```
while b != 0:
    carry = a AND b           (positions that carry)
    a = a XOR b               (sum without carry)
    b = carry << 1           (carry moves one bit left)
return a
```

```

// VARIATION 1: mod-3 state machine (Single Number II)
// MENTAL MODEL: extend XOR from mod-2 (cancel pairs) to mod-3 (cancel triples);
//                ones/twos track bits seen 1 or 2 times mod 3.
// Instead of XOR canceling pairs (mod 2), cancel triples (mod 3)
// ones = bits seen exactly 1 mod 3 times; twos = bits seen exactly 2 mod 3 times
int ones = 0, twos = 0;
for (int num : nums) {
    ones = (ones ^ num) & ~twos;    // ← add to ones if not already in twos
    twos = (twos ^ num) & ~ones;    // ← add to twos if not already in ones
}
// After all: ones holds the element appearing once (0 mod 3 remainder = survive)

// VARIATION 2: split XOR into two groups (Single Number III)
// Two unique elements a, b. XOR of all = a ^ b.
// Use lowest set bit of (a^b) to split nums into two groups → each group has one unique.
int xor = 0;
for (int num : nums) xor ^= num;    // xor = a ^ b
int diff = xor & (-xor);            // ← lowest set bit that differs between a and b
int a = 0;
for (int num : nums)
    if ((num & diff) != 0) a ^= num; // ← XOR one group → one unique element
int b = xor ^ a;                    // ← the other unique element

// VARIATION 3: DP relation (Counting Bits)
// dp[i] = dp[i/2] + last bit of i
// i >> 1 = i without last bit (same or fewer ones); last bit = i & 1
int[] dp = new int[n + 1];
for (int i = 1; i <= n; i++)
    dp[i] = dp[i >> 1] + (i & 1);    // ← right-shift halves the problem

// VARIATION 4: common prefix (Bitwise AND of Range)
// AND of any range [left, right]: bits that differ between left and right get cleared.
// Right-shift both until equal → find shared high-bit prefix → shift back.
int shift = 0;
while (left != right) { left >>= 1; right >>= 1; shift++; }
return left << shift;               // ← shift back to original magnitude

// VARIATION 5: carry loop (Sum of Two Integers)
// XOR gives bit sum without carry; AND shifted left gives carry.
// Repeat until no carry remains.
int add(int a, int b) {
    while (b != 0) {
        int carry = a & b;          // ← carry bits
        a = a ^ b;                  // ← partial sum (no carry)
        b = carry << 1;             // ← carry propagated one position left
    }
    return a;
}

```

Part 1 — XOR Cancel

#136 Single Number

Description: Array where every element appears twice except one. Find it.

Intuition: XOR is its own inverse, so every duplicated pair cancels to 0 and only the lone element is left standing.

Algorithm: XOR all — every pair cancels ($a \oplus a = 0$), leaving the single element.

Variables: `result` = XOR accumulator — paired values cancel, the lone value survives · `num` = current element

Pseudocode:

```
result = 0
for each num in nums:
    result = result XOR num    (a XOR a = 0, so duplicate pairs vanish)
return result                (only the single element remains)
```

```
class Solution {
    public int singleNumber(int[] nums) {
        int result = 0;
        for (int num : nums) result ^= num;    // pairs cancel; lone element survives
        return result;
    }
}
```

Time $O(n)$ | **Space** $O(1)$

#268 Missing Number

Description: Array of n distinct numbers from $[0, n]$ with one missing. Find it.

Intuition: XOR-ing every index against every value pairs each present number with its index — the missing number's index has no partner to cancel.

Variation: XOR each index `i` with `nums[i]`. Indices $0..n$ XOR'd with the n elements — the missing index has no pair.

Variables: `result` = XOR accumulator seeded with n (the absent last index) · `i` = index · `nums[i]` = value at that index

Pseudocode:

```
result = n                (start with the largest index, which has no value slot)
for i = 0..n-1:
    result = result XOR i XOR nums[i]    (each present value cancels its index)
return result              (only the missing number's index has no partner)
```

```
class Solution {
    public int missingNumber(int[] nums) {
        int result = nums.length;                // start with n (last index)
        for (int i = 0; i < nums.length; i++)
            result ^= i ^ nums[i];                // ← VARIATION: XOR index with value
        return result;
    }
}
```

Time $O(n)$ | **Space** $O(1)$

#137 Single Number II

Description: Every element appears three times except one. Find it.

Intuition: XOR cancels pairs (mod 2); here we need a counter mod 3, so two state bits (`ones`, `twos`) cycle each bit through $1 \rightarrow 2 \rightarrow 0$ and the survivor sits in `ones`.

Variation: Two-state XOR machine. `ones` tracks bits seen 1 mod 3 times, `twos` tracks bits seen 2 mod 3 times. When a bit reaches 3, it clears from both.

Variables: `ones` = bits seen 1 time mod 3 · `twos` = bits seen 2 times mod 3 · `num` = current element (a bit reaching 3 occurrences clears from both) **Pseudocode:**

```

ones = 0; twos = 0
for each num in nums:
    ones = (ones XOR num) AND NOT twos    (add bit to ones unless it is already in twos)
    twos = (twos XOR num) AND NOT ones    (add bit to twos unless it is now in ones)
return ones                               (bits at count 0 mod 3 survive here = the single element)

```

```

class Solution {
    public int singleNumber(int[] nums) {
        int ones = 0, twos = 0;
        for (int num : nums) {
            ones = (ones ^ num) & ~twos;    // ← VARIATION: add to ones only if not in twos
            twos = (twos ^ num) & ~ones;    // ← VARIATION: add to twos only if not in ones
        }
        return ones;                       // ones holds the element seen exactly once
    }
}

```

Time $O(n)$ | **Space** $O(1)$

#260 Single Number III

Description: Two elements appear exactly once; all others appear twice. Find both.

Intuition: XOR of all leaves $a \oplus b$; any set bit in it is a bit where a and b differ, so it splits the array into two groups each holding one unique.

Variation: XOR all $\rightarrow$ get $a \oplus b$. Use lowest set bit of $a \oplus b$ to partition nums into two groups (one containing a , one containing b). XOR each group independently.

Variables: xor = XOR of all = $a \oplus b$ · $diff$ = lowest set bit of xor (a bit where a and b differ) · a = unique value isolated from the group whose $diff$ bit is set · num = current element **Pseudocode:**

```

xor = 0
for each num: xor = xor XOR num            (everything else cancels -> xor = a XOR b)
diff = xor AND (-xor)                     (isolate one bit where a and b differ)
a = 0
for each num:
    if (num AND diff) != 0: a = a XOR num (XOR only the group with that bit set)
return [a, xor XOR a]                     (the other unique is xor XOR a)

```

```

class Solution {
    public int[] singleNumber(int[] nums) {
        int xor = 0;
        for (int num : nums) xor ^= num;    // xor = a ^ b

        int diff = xor & (-xor);            // ← VARIATION: isolate lowest bit that differs
        int a = 0;
        for (int num : nums)
            if ((num & diff) != 0) a ^= num; // ← VARIATION: XOR only one group
        return new int[]{a, xor ^ a};       // b = (a^b) ^ a
    }
}

```

Time $O(n)$ | **Space** $O(1)$

Part 2 — Bit Counting

#191 Number of 1 Bits

Description: Return the number of set bits in an unsigned integer.

Intuition: Each $n \& (n-1)$ erases exactly one set bit, so the loop runs once per set bit — no need to scan all 32 positions.

Algorithm: Brian Kernighan — $n \& (n-1)$ clears the lowest set bit each iteration.

Variables: n = number being whittled down · $count$ = set bits cleared so far ($n \& (n-1)$ removes the lowest set bit each step) **Pseudocode:**

```
count = 0
while n != 0:
    count = count + 1
    n = n AND (n - 1)      (clears the lowest set bit)
return count              (loop ran once per set bit)
```

```
public class Solution {
    public int hammingWeight(int n) {
        int count = 0;
        while (n != 0) {
            count++;
            n &= (n - 1);    // ← clears lowest set bit; loop runs exactly hamming-weight times
        }
        return count;
    }
}
```

Time $O(k)$ k = number of set bits | **Space** $O(1)$

#338 Counting Bits

Description: Return an array `result` where `result[i]` = number of 1 bits in `i`, for i in $[0, n]$.

Intuition: i has the same set bits as $i \gg 1$ plus possibly its own last bit, so reuse the already-computed smaller answer instead of recounting.

Variation: DP relation — $i \gg 1$ is i with last bit removed (already computed), plus the last bit $i \& 1$.

Variables: `dp[]` = set-bit count for each number $0..n$ · i = current number ($i \gg 1$ is the already-solved smaller number, $i \& 1$ is its last bit) **Pseudocode:**

```
dp = array of size n+1 (dp[0] = 0)
for i = 1..n:
    dp[i] = dp[i >> 1] + (i AND 1)    (count of i without its last bit, plus that bit)
return dp
```

```
class Solution {
    public int[] countBits(int n) {
        int[] dp = new int[n + 1];
        for (int i = 1; i <= n; i++)
            dp[i] = dp[i >> 1] + (i & 1);    // ← VARIATION: dp[i/2] + last bit
        return dp;
    }
}
```

Time $O(n)$ | Space $O(n)$

#190 Reverse Bits

Description: Reverse the 32 bits of an unsigned integer.

Intuition: Peel the lowest bit off `n` and stack it onto `result` from the bottom — after 32 shifts the bit order is fully mirrored.

Variation: Fixed 32 iterations — each time take the LSB of `n`, append it to `result`, then shift both.

Variables: `result` = mirrored bits being built from the bottom up · `n` = source bits being peeled from the bottom · `i` = iteration counter (always 32) **Pseudocode:**

```
result = 0
repeat 32 times (i = 0..31):
    result = (result << 1) OR (n AND 1)    (shift result up, drop in n's lowest bit)
    n = n >> 1                            (advance to next source bit)
return result
```

```
public class Solution {
    public int reverseBits(int n) {
        int result = 0;
        for (int i = 0; i < 32; i++) {      // ← VARIATION: always 32 iterations (fixed-width)
            result = (result << 1) | (n & 1); // append LSB of n to result
            n >>= 1;
        }
        return result;
    }
}
```

Time $O(32) = O(1)$ | Space $O(1)$

Part 3 — Range / Arithmetic

#201 Bitwise AND of Numbers Range

Description: Return the AND of all numbers in `[left, right]`.

Intuition: AND only keeps bits that are 1 in every number; the low bits flip somewhere in the range, so only the common high-bit prefix of `left` and `right` survives.

Key insight: any two adjacent numbers in the range differ in at least the lowest bit. AND of a range clears all bits that differ across any pair in the range. The result is the common high-bit prefix of `left` and `right`.

Variation: right-shift both until equal to find shared prefix; shift back.

Variables: `left / right` = range endpoints, shifted right until they match · `shift` = number of low bits dropped (restored at the end) **Pseudocode:**

```
shift = 0
while left != right:                (low bits differ -> they get cleared by AND)
    left = left >> 1
    right = right >> 1
    shift = shift + 1
return left << shift                (shared high-bit prefix, shifted back into place)
```

```

class Solution {
    public int rangeBitwiseAnd(int left, int right) {
        int shift = 0;
        while (left != right) {          // ← VARIATION: shift until common prefix found
            left  >>= 1;
            right >>= 1;
            shift++;
        }
        return left << shift;           // ← restore to original magnitude
    }
}

```

Time $O(\log n)$ | **Space** $O(1)$

#371 Sum of Two Integers

Description: Return $a + b$ without using $+$ or $-$.

Intuition: Addition splits into a carry-free sum (XOR) and the carries (AND shifted left); feed the carry back until nothing carries over.

Variation: XOR computes bit sum without carry; AND shifted left computes carry. Repeat until no carry.

Variables: a = running partial sum (no carry) · b = pending carry to add next · $carry$ = bit positions where both currently have 1 **Pseudocode:**

```

while b != 0:                (still have carry to fold in)
    carry = a AND b          (positions that generate a carry)
    a = a XOR b              (sum of the bits ignoring carry)
    b = carry << 1           (carry takes effect one position to the left)
return a                     (no carry left → a is the sum)

```

```

class Solution {
    public int getSum(int a, int b) {
        while (b != 0) {
            int carry = a & b;          // ← VARIATION: carry = bits where both are 1
            a = a ^ b;                  // ← VARIATION: partial sum (no carry propagation)
            b = carry << 1;             // ← VARIATION: carry propagates one bit left
        }
        return a;
    }
}

```

Time $O(32) = O(1)$ | **Space** $O(1)$

Part 4 — Bitmask as Set

#318 Maximum Product of Word Lengths

Description: Given a list of words, find the maximum product $\text{len}(a) * \text{len}(b)$ where a and b share no common letters.

Intuition: Compress each word's letter set into a 26-bit integer so "share a letter?" becomes a single AND, replacing slow character-by-character comparison.

Variation: encode each word as a 26-bit integer where bit $i = 1$ if letter $'a' + i$ appears. Two words share a letter iff their bitmasks have any common bit ($\text{mask}[i] \& \text{mask}[j] \neq 0$).

Variables: `masks[]` = 26-bit letter-set per word (bit `i` set if letter 'a'+`i` appears) · `i / j` = word indices being compared ·
`max` = best product of lengths so far (disjoint sets have `AND == 0`) **Pseudocode:**

```
build masks: for each word i, for each char c, set bit (c - 'a') in masks[i]
max = 0
for i = 0..n-1:
    for j = i+1..n-1:
        if (masks[i] AND masks[j]) == 0:            (no shared letter)
            max = larger of (max, len(word i) * len(word j))
return max
```

```
class Solution {
    public int maxProduct(String[] words) {
        int n = words.length;
        int[] masks = new int[n];
        for (int i = 0; i < n; i++)
            for (char c : words[i].toCharArray())
                masks[i] |= (1 << (c - 'a')); // ← VARIATION: 26-bit bitmask per word

        int max = 0;
        for (int i = 0; i < n; i++)
            for (int j = i + 1; j < n; j++)
                if ((masks[i] & masks[j]) == 0) // ← VARIATION: no shared letters = AND is zero
                    max = Math.max(max, words[i].length() * words[j].length());
        return max;
    }
}
```

Time $O(n^2 + n \cdot k)$ k = avg word length | **Space** $O(n)$

Bit Trick Cheat Sheet

```
n & (n-1)           // clear lowest set bit → count bits, check power of 2
n & (-n)            // isolate lowest set bit → split into groups
n ^ n = 0           // same value cancels → XOR pairs
n ^ 0 = n           // identity → XOR with index
(n >> i) & 1        // check bit i
n | (1 << i)        // set bit i
n & ~(1 << i)       // clear bit i
n ^ (1 << i)        // toggle bit i
i >> 1              // divide by 2 (fast)
i & 1               // last bit (0 = even, 1 = odd)

// --- Divisibility by powers of two ---
(n & 1) == 0        // divisible by 2 (even) ← LSB is the 2^0 place
(n & 1) == 1        // NOT divisible by 2 (odd)
(n & ((1 << k) - 1)) // remainder of n / 2^k ; == 0 means divisible by 2^k
(n & 3) == 0        // divisible by 4 (k = 2)
(n & 7) == 0        // divisible by 8 (k = 3)
// Note: n & 1 is safe for NEGATIVE numbers (two's complement) where
//         n % 2 returns -1 for odd negatives. Prefer & 1 for parity.
```

Pattern Chooser

Signal	Technique	Time
"appears odd/once, rest appear even/twice"	XOR all	$O(n)$
"appears once, rest appear 3 times"	mod-3 state machine (ones , twos)	$O(n)$
"two unique elements"	XOR $\rightarrow$ split by <code>xor & (-xor)</code>	$O(n)$
Count set bits	Brian Kernighan <code>n &= (n-1)</code>	$O(k)$ k = set bits
Count bits for all $0..n$	DP: <code>dp[i] = dp[i>>1] + (i&1)</code>	$O(n)$
AND of a range	Right-shift to find common prefix	$O(\log n)$
Add without +	XOR + carry loop	$O(1)$ (32 iters)
"no shared characters between two strings"	26-bit bitmask, check <code>AND == 0</code>	$O(n^2)$ pairwise
"is n a power of two?"	<code>n > 0 && (n & (n-1)) == 0</code>	$O(1)$
"is n a power of four?"	Power of 2 AND set bit at even position: <code>(n & 0xAAAAAAAA) == 0</code>	$O(1)$
"find duplicate, no extra space, no modification"	Floyd's cycle detection on array-as-linked-list	$O(n)$

Part 5 — Power Checks & Cycle Detection

#231 Power of Two

Description: Given an integer `n`, return true if it is a power of two.

Intuition: A power of two is a single 1 bit, and `n & (n-1)` clears that one bit to 0 — anything else leaves bits behind.

Algorithm: A power of two has exactly one set bit. Check `n > 0` and `n & (n-1) == 0` (clearing the lowest set bit results in 0 only if there was one bit).

Variables: `n` = number tested (`n & (n-1)` clears its single set bit, leaving 0 only when exactly one bit was set)

Pseudocode:

```
return true when:
    n > 0                (must be positive)
    AND (n AND (n-1)) == 0 (clearing the lone set bit leaves nothing)
```

```
class Solution {
    public boolean isPowerOfTwo(int n) {
        return n > 0 && (n & (n - 1)) == 0; // ← VARIATION: single-bit property of powers of 2
    }
}
```

Time $O(1)$ | **Space** $O(1)$

#342 Power of Four

Description: Given an integer `n`, return true if it is a power of four.

Intuition: Powers of four are powers of two (one set bit) whose bit sits at an even position; masking against the odd-position mask `0xAAAAAAAA` must give 0.

Variation: Power of four must be power of two (one set bit) AND that bit must be at an even bit position (0, 2, 4, 6, ...). Mask `0xAAAAAAAA` has bits set at ALL odd positions; if `n & 0xAAAAAAAA == 0`, the set bit is at an even position.

Variables: `n` = number tested · mask `0xAAAAAAAA` = all odd bit positions (a power of four's lone bit must sit at an even position, so AND with this mask is 0) **Pseudocode:**

```
return true when all hold:
    n > 0
    (n AND (n-1)) == 0          (exactly one set bit -> power of two)
    (n AND 0xAAAAAAAA) == 0    (that bit is at an even position -> power of four)
```

```
class Solution {
    public boolean isPowerOfFour(int n) {
        // 0xAAAAAAAA = 1010 1010 ... 1010 (bits at odd positions 1,3,5,...31)
        return n > 0 && (n & (n-1)) == 0          // ← power of 2: exactly one set bit
            && (n & 0xAAAAAAAA) == 0;           // ← VARIATION: set bit at even position (4^k has bit
    }
}
```

Time $O(1)$ | **Space** $O(1)$

#287 Find the Duplicate Number

Description: Given an array of `n+1` integers where each is in `[1, n]`, find the duplicate. No extra space, no modifying the array.

Intuition: Following `i → nums[i]` turns the array into a linked list; a duplicate value means two indices point to the same node, forming a cycle whose entry is the duplicate.

Variation: Floyd's Cycle Detection. Treat the array as a linked list where `nums[i]` is the next node. Since there's a duplicate, two indices point to the same value → creates a cycle. Find cycle entry = duplicate.

Variables: `slow` = single-step pointer · `fast` = double-step pointer (`nums[i]` acts as the "next" link; the cycle entry is the duplicate) **Pseudocode:**

```
Phase 1 – find a meeting point inside the cycle:
    slow = nums[0]; fast = nums[0]
    do: slow = nums[slow]; fast = nums[nums[fast]]    until slow == fast

Phase 2 – find the cycle entrance:
    slow = nums[0]
    while slow != fast: slow = nums[slow]; fast = nums[fast]
    return slow      (cycle entry = the duplicate value)
```

```

class Solution {
    public int findDuplicate(int[] nums) {
        // Phase 1: find intersection point in cycle
        int slow = nums[0], fast = nums[0];
        do {
            slow = nums[slow];           // ← VARIATION: Floyd's cycle, not XOR
            fast = nums[nums[fast]];
        } while (slow != fast);
        // Phase 2: find cycle entry (= duplicate)
        slow = nums[0];
        while (slow != fast) {
            slow = nums[slow];
            fast = nums[fast];
        }
        return slow;
    }
}

```

Time $O(n)$ | **Space** $O(1)$

Additional Reference Problems

#405 Convert a Number to Hexadecimal

Description: Given an integer `num`, return its hexadecimal representation as a lowercase string. For negative numbers use two's complement.

Intuition: Hex is base 16, so each group of 4 bits maps to one hex digit; mask the lowest 4 bits with `num & 15` and shift right 4 at a time, using an unsigned shift so the two's-complement sign bit fills with zeros.

Algorithm: Repeatedly take `num & 0xF` to read the lowest nibble, map it to a hex char, then unsigned-right-shift `num >> 4`. Stop when `num` becomes 0.

Variables: `num` = number being consumed 4 bits at a time · `digits` = hex digit lookup table · `nibble` = lowest 4 bits (`num & 15`) = one hex digit · `builder` = digits collected low-to-high (reversed at the end) **Pseudocode:**

```

if num == 0: return "0"
digits = "0123456789abcdef"
builder = empty
while num != 0:
    nibble = num AND 15           (lowest 4 bits = one hex digit)
    append digits[nibble] to builder
    num = num >> 4                (unsigned shift fills sign bit with 0)
return builder reversed          (we appended least-significant digit first)

```

```

class Solution {
    public String toHex(int num) {
        if (num == 0) {
            return "0";
        }
        char[] digits = "0123456789abcdef".toCharArray();
        StringBuilder builder = new StringBuilder();
        while (num != 0) {
            int nibble = num & 15; // ← lowest 4 bits = one hex digit
            builder.append(digits[nibble]);
            num >>= 4; // unsigned shift handles negatives (two's complement)
        }
        return builder.reverse().toString();
    }
}

```

Time $O(8) = O(1)$ | **Space** $O(1)$

#461 Hamming Distance

Description: Given two integers x and y , return the Hamming distance — the number of bit positions where they differ.

Intuition: XOR sets exactly the bits where x and y differ, so the answer is just the number of set bits in $x \oplus y$.

Algorithm: Compute $x \oplus y$, then count its set bits with Brian Kernighan ($n \& (n-1)$ clears the lowest set bit each iteration).

Variables: $n = x \oplus y$ (1 exactly where the two differ) · count = differing bits counted ($n \& (n-1)$ clears one set bit per step) **Pseudocode:**

```

n = x XOR y           (set bits mark positions where x and y differ)
count = 0
while n != 0:
    count = count + 1
    n = n AND (n - 1)  (clear lowest set bit)
return count

```

```

class Solution {
    public int hammingDistance(int x, int y) {
        int n = x ^ y; // differing bits become 1
        int count = 0;
        while (n != 0) {
            count++;
            n &= (n - 1); // ← clears lowest set bit each iteration
        }
        return count;
    }
}

```

Time $O(k)$ k = differing bits | **Space** $O(1)$

#476 Number Complement

Description: Given a positive integer num , return its complement: flip every bit up to and including the most significant set bit.

Intuition: Build a mask of all 1s spanning the bit width of `num` (from the MSB down to bit 0), then XOR — flipping exactly the meaningful bits while leaving the leading zeros untouched.

Algorithm: Grow a mask `1 << i` until it covers all bits of `num` (`mask < num`), forming `(1 << bitLength) - 1`; XOR `num` with that mask to flip only the relevant bits.

Variables: `num` = input whose meaningful bits get flipped · `mask` = all-ones mask spanning `num`'s bit width · `i` = bit position being added to the mask **Pseudocode:**

```
mask = 0; i = 0
while (1 << i) <= num:           (extend mask up to num's most significant bit)
    set bit i in mask
    i = i + 1
return num XOR mask              (flips every bit within num's width, leaves leading zeros)
```

```
class Solution {
    public int findComplement(int num) {
        int mask = 0;
        int i = 0;
        while ((1 << i) <= num) {           // ← extend mask to num's bit width
            mask |= (1 << i);               // set bit i in the all-ones mask
            i++;
        }
        return num ^ mask;                  // XOR flips every bit within the mask
    }
}
```

Time $O(\log \text{num})$ | **Space** $O(1)$

#957 Prison Cells After N Days

Description: 8 prison cells (`cells[i]` is 0 or 1) update each day: a cell becomes 1 if both neighbors are equal, else 0. The two end cells always become 0 (they lack two neighbors). Return the state after `n` days.

Intuition: With only 8 cells the state space is finite, so the configuration must cycle; detect the cycle length and reduce `n` modulo it to avoid simulating huge day counts.

Algorithm: Encode the 8 cells as an 8-bit integer. Simulate one day with bit operations (a cell is 1 iff its neighbors match, i.e. their XOR is 0), recording seen states in a map. On the first repeat, reduce remaining days modulo the cycle length, then finish the leftover days.

Variables: `state` = 8 cells packed into an 8-bit integer · `n` = remaining days · `seen` = map from state → the `n` value when first seen (to measure cycle length) · `result[]` = unpacked cells · `nextDay` helper uses `left / right` neighbor bits, setting cell `i` when they are equal **Pseudocode:**

```

pack cells into state: for i = 0..7, if cells[i]==1 set bit i of state
seen = empty map
while n > 0:
    if state already in seen:
        n = n mod (seen[state] - n)          (skip whole repeated cycles)
    else:
        seen[state] = n
    if n > 0:
        state = nextDay(state)
        n = n - 1
unpack: for i = 0..7, result[i] = bit i of state
return result

nextDay(state):
    next = 0
    for i = 1..6:                          (end cells 0 and 7 always become 0)
        left  = bit (i-1) of state
        right = bit (i+1) of state
        if (left XOR right) == 0: set bit i of next    (neighbors equal)
    return next

```

```

import java.util.HashMap;
import java.util.Map;

class Solution {
    public int[] prisonAfterNDays(int[] cells, int n) {
        int state = 0;
        for (int i = 0; i < 8; i++) {
            if (cells[i] == 1) {
                state |= (1 << i);          // pack cells into an 8-bit integer
            }
        }
        Map<Integer, Integer> seen = new HashMap<>();
        while (n > 0) {
            if (seen.containsKey(state)) {
                n %= seen.get(state) - n;    // ← cycle found: skip whole cycles
            } else {
                seen.put(state, n);
            }
            if (n > 0) {
                state = nextDay(state);
                n--;
            }
        }
        int[] result = new int[8];
        for (int i = 0; i < 8; i++) {
            result[i] = (state >> i) & 1;    // unpack bits back into cells
        }
        return result;
    }

    private int nextDay(int state) {
        int next = 0;
        for (int i = 1; i < 7; i++) {
            // end cells (0 and 7) always become 0
            int left = (state >> (i - 1)) & 1;
            int right = (state >> (i + 1)) & 1;
            if ((left ^ right) == 0) {
                // neighbors equal → cell becomes 1
                next |= (1 << i);
            }
        }
        return next;
    }
}

```

Time $O(\min(n, \text{cycle}) \cdot 8)$ | **Space** $O(\text{states})$